

austin_race_analysis

May 13, 2025

0.1 Unsupervised Learning to Predict Race from Merged Dataset

```
[1]: #imports
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

0.2 Load Data and Inspect

```
[2]: full_data_df = pd.read_csv('final_merged.csv')
full_data_df

print("Shape of the data: ", full_data_df.shape)

full_data_df.info()
full_data_df.head()
```

Shape of the data: (9222, 140)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9222 entries, 0 to 9221

Columns: 140 entries, Patient ID to Survival Months

dtypes: float64(77), int64(2), object(61)

memory usage: 9.9+ MB

/tmp/ipykernel_29693/3584519652.py:1: DtypeWarning: Columns (0,3,4,5,6,7,9,12,15,16,17,18,22,24,26,28,29,30,40,65,66,67,68,69,70,71,72,73,74,75,76,82,83,89,90,91,92,93,96,97,98,99,100,101,102,103,104,105,106,107,108,109,115) have mixed types. Specify dtype option on import or set low_memory=False.

```
full_data_df = pd.read_csv('final_merged.csv')
```

```
[2]:
```

	Patient ID	Age	Surgery	Cancer Type	Cancer Type	Detailed	Cellularity	\
0	Breast_MRI_001	41	NaN	NaN		NaN	NaN	
1	Breast_MRI_001	41	NaN	NaN		NaN	NaN	
2	Breast_MRI_002	38	NaN	NaN		NaN	NaN	
3	Breast_MRI_003	62	NaN	NaN		NaN	NaN	
4	Breast_MRI_003	62	NaN	NaN		NaN	NaN	

	Chemotherapy	Pam50 + Claudin-low subtype	Cohort	ER status measured by IHC	\
0	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	

	...	Unnamed: 3	T Stage	6th Stage	Grade	\
0	...	NaN	T2	IIB	Poorly differentiated; Grade III	
1	...	NaN	T2	IIB	Moderately differentiated; Grade II	
2	...	NaN	NaN	NaN	NaN	
3	...	NaN	T2	IIB	Poorly differentiated; Grade III	
4	...	NaN	T2	IIB	Moderately differentiated; Grade II	

	A Stage	Estrogen Status	Progesterone Status	Regional Node Examined	\
0	Regional	Positive	Negative	6.0	
1	Regional	Positive	Positive	9.0	
2	NaN	NaN	NaN	NaN	
3	Regional	Positive	Positive	6.0	
4	Regional	Positive	Positive	19.0	

	Reginol Node Positive	Survival Months
0	1.0	66.0
1	1.0	72.0
2	NaN	NaN
3	1.0	52.0
4	1.0	55.0

[5 rows x 140 columns]

0.3 Data Cleaning

```
[13]: #Dropping any duplicate rows
full_data_df.drop_duplicates(inplace=True)

# Drop features with a lot of missing values
missing__ratio = full_data_df.isnull().mean()
columns_to_drop = missing__ratio[missing__ratio > 0.8].index
full_data_df.drop(columns=columns_to_drop, inplace=True)

[15]: # Look at new shape of the data
print("Shape of the data after dropping columns with a lot of missing values:␣
↵", full_data_df.shape)
```

Shape of the data after dropping columns with a lot of missing values: (6693, 111)

```
[16]: #inspecting the race column
full_data_df['Race'].value_counts(dropna=False)

#Drop rows with missing values in Race column

full_data_df = full_data_df.dropna(subset=['Race'])

#Drop rows where race is = 0
full_data_df = full_data_df[full_data_df['Race'] != 0]

full_data_df['Race'].value_counts()
```

```
[16]: Race
1.0    5856
2.0     470
4.0    321
5.0     18
3.0     14
6.0      9
8.0      4
7.0      1
Name: count, dtype: int64
```

0.4 EDA

```
[17]: full_data_df.describe()
```

```
[17]:
```

	Age	ER	HER2	Histologic type	PR \
count	6693.000000	4944.000000	4944.000000	6693.000000	4944.000000
mean	53.247124	0.803600	0.189523	0.596145	0.702063
std	8.718803	0.397314	0.391963	0.623003	0.457398
min	21.000000	0.000000	0.000000	0.000000	0.000000
25%	47.000000	1.000000	0.000000	0.000000	0.000000
50%	53.000000	1.000000	0.000000	1.000000	1.000000
75%	60.000000	1.000000	0.000000	1.000000	1.000000
max	89.000000	1.000000	1.000000	9.000000	1.000000

	Tumor Size	T Stage	Days to MRI (From the Date of Diagnosis) \
count	6056.000000	6687.000000	4944.000000
mean	28.519815	1.708689	12.528519
std	19.264544	0.714212	17.320590
min	1.000000	1.000000	-123.000000
25%	16.000000	1.000000	7.000000
50%	23.000000	2.000000	12.000000
75%	35.000000	2.000000	19.250000
max	140.000000	4.000000	381.000000

	Manufacturer	Manufacturer Model Name	...	\
count	4944.000000	4944.000000	...	
mean	0.582120	2.893811	...	
std	0.908594	1.904946	...	
min	0.000000	0.000000	...	
25%	0.000000	1.000000	...	
50%	0.000000	3.000000	...	
75%	2.000000	4.000000	...	
max	2.000000	7.000000	...	

Pathologic response to Neoadjuvant therapy: Pathologic stage (N)
following neoadjuvant therapy \

count	1722.000000
mean	0.660859
std	0.837785
min	-1.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	3.000000

Pathologic response to Neoadjuvant therapy: Pathologic stage (M)
following neoadjuvant therapy \

count	1673.000000
mean	-0.982666
std	0.180538
min	-1.000000
25%	-1.000000
50%	-1.000000
75%	-1.000000
max	1.000000

Overall Near-complete Response: Stricter Definition \

count	1929.000000
mean	0.444272
std	0.770387
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000

Overall Near-complete Response: Looser Definition \

count	1929.000000
mean	0.670295
std	0.890914
min	0.000000

25%	0.000000
50%	0.000000
75%	1.000000
max	3.000000


	Near-complete Response (Graded Measure)	Race	N Stage \
count	1929.00000	6693.000000	6671.000000
mean	2.39606	1.240849	1.204317
std	4.75256	0.750747	0.685523
min	0.00000	1.000000	0.000000
25%	0.00000	1.000000	1.000000
50%	0.00000	1.000000	1.000000
75%	1.00000	1.000000	1.000000
max	19.00000	8.000000	3.000000

	Regional Node Examined	Reginol Node Positive	Survival Months
count	6056.000000	6056.000000	6056.000000
mean	13.934610	3.410337	72.084544
std	8.070807	4.475298	22.470205
min	1.000000	1.000000	1.000000
25%	8.000000	1.000000	57.000000
50%	13.000000	2.000000	74.000000
75%	18.000000	3.000000	90.000000
max	61.000000	46.000000	107.000000

[8 rows x 67 columns]

```
[18]: #Coorelation heatmap for numerical features

numeric_features = full_data_df.select_dtypes(include=[np.number])
plt.figure(figsize=(12, 8))
sns.heatmap(numeric_features.corr(), annot=False, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



0.5 Guassian and PCA

[46]: `from sklearn.preprocessing import StandardScaler`

```
# Assuming 'full_data_df' is your merged breast cancer dataset
# Define race label and drop it from feature set
target_col = 'Race'

# 1. Drop columns with too much missing data (>80% missing already dropped earlier)
# 2. Keep only numeric columns (Naive Bayes needs numeric features)
numeric_cols = full_data_df.select_dtypes(include=[np.number]).columns.
drop(target_col)
features = numeric_cols

# Strategy 1: Drop rows with any missing values
```

```

df_dropped = full_data_df[features.union([target_col])].dropna()

# Strategy 2: Fill numeric NaNs with column means
df_filled = full_data_df[features.union([target_col])].copy()
df_filled[features] = df_filled[features].fillna(df_filled[features].mean())

# Scale features
scaler = StandardScaler()
df_dropped[features] = scaler.fit_transform(df_dropped[features])
df_filled[features] = scaler.fit_transform(df_filled[features])

print("Dropped dataset shape:", df_dropped.shape)
print("Filled dataset shape:", df_filled.shape)

```

Dropped dataset shape: (118, 67)
 Filled dataset shape: (6693, 67)

```

[48]: from sklearn.model_selection import train_test_split
      from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import accuracy_score, confusion_matrix
      import matplotlib.pyplot as plt
      import seaborn as sns
      import random

# Define features and target
X = df_filled[features]
y = df_filled[target_col]

# Drop classes with fewer than 2 samples
class_counts = y.value_counts()
valid_classes = class_counts[class_counts >= 2].index

X = X[y.isin(valid_classes)]
y = y[y.isin(valid_classes)]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
    ↳test_size=0.2, random_state=42)

# Test different subset sizes
feature_sizes = [5, 7, 10]
results = {}

for size in feature_sizes:
    selected_features = random.sample(list(X.columns), size)

```

```

model = GaussianNB()
model.fit(X_train[selected_features], y_train)
y_pred = model.predict(X_test[selected_features])

acc = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

results[size] = {
    "features": selected_features,
    "accuracy": acc,
    "conf_matrix": conf_matrix
}

print(f"\nFeature Subset Size: {size}")
print("Selected Features:", selected_features)
print(f"Accuracy: {acc:.4f}")


plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title(f"Confusion Matrix (Feature Size: {size})")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

Feature Subset Size: 5

Selected Features: ['Tumor Size', 'Recurrence event(s)', 'Pathologic Response to Neoadjuvant Therapy', 'Oncotype score', 'Staging(Metastasis)#(Mx -replaced by -1) [M] ']


Accuracy: 0.6490



Feature Subset Size: 7

Selected Features: ['Manufacturer Model Name', 'Age', 'Overall Near-complete Response: Looser Definition', 'Adjuvant Chemotherapy', 'Reginol Node Positive', 'Mol Subtype', 'N Stage']


Accuracy: 0.6617



Feature Subset Size: 10

Selected Features: ['Age at last contact in EMR f/u(days)(from the date of diagnosis) ,last time patient known to be alive, unless age of death is reported(in such case the age of death', 'Overall Near-complete Response: Stricter Definition', 'Neoadjuvant Radiation Therapy', 'Reconstruction Diameter ', 'Skin/Nipple Invovlement', 'PR', 'Reginol Node Positive', 'Known Ovarian Status ', 'FOV Computed (Field of View) in cm ', 'Flip Angle ']

Accuracy: 0.6580



```
[49]: # Baseline model
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Train full-feature Naive Bayes model
model_full = GaussianNB()
model_full.fit(X_train, y_train)

# Predict on test set
y_pred_full = model_full.predict(X_test)

# Evaluate
accuracy_full = accuracy_score(y_test, y_pred_full)
print(f"Naive Bayes Accuracy (All Features): {accuracy_full:.4f}")

# Classification report
print("Classification Report:\n", classification_report(y_test, y_pred_full))

# Confusion Matrix
conf_matrix_full = confusion_matrix(y_test, y_pred_full)
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_full, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix (All Features)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

Naive Bayes Accuracy (All Features): 0.7588

Classification Report:

	precision	recall	f1-score	support
1.0	1.00	0.76	0.86	1172
2.0	0.67	0.66	0.67	94
3.0	1.00	1.00	1.00	3
4.0	0.17	0.97	0.30	64
5.0	1.00	1.00	1.00	3
6.0	1.00	0.50	0.67	2
8.0	0.00	0.00	0.00	1
accuracy			0.76	1339
macro avg	0.69	0.70	0.64	1339
weighted avg	0.94	0.76	0.82	1339

/home/austi/miniconda3/envs/msse/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/home/austi/miniconda3/envs/msse/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/home/austi/miniconda3/envs/msse/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.


_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))



```
[50]: from sklearn.decomposition import PCA

# Reduce to 2D for visualization
pca_vis = PCA(n_components=2)
X_train_pca_2d = pca_vis.fit_transform(X_train)

# Plot PCA components colored by Race
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_train_pca_2d[:, 0], y=X_train_pca_2d[:, 1], hue=y_train,
               palette="Set1")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.title("2D PCA Projection Colored by Race")
plt.legend(title="Race")
plt.grid(True)
plt.tight_layout()
plt.show()
```



0.6 Evaluation

```
[51]: from sklearn.decomposition import PCA
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report

# Fit PCA to retain 90% variance
pca_final = PCA(n_components=0.90)
X_train_pca = pca_final.fit_transform(X_train)
X_test_pca = pca_final.transform(X_test)

print(f"Number of PCA components selected: {pca_final.n_components}")

# Train and predict
model_pca = GaussianNB()
model_pca.fit(X_train_pca, y_train)
y_pred_pca = model_pca.predict(X_test_pca)

# Evaluate
```

```

acc_pca = accuracy_score(y_test, y_pred_pca)
print(f"Naive Bayes Accuracy (PCA Features): {acc_pca:.4f}")


conf_matrix_pca = confusion_matrix(y_test, y_pred_pca)
sns.heatmap(conf_matrix_pca, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix (PCA Features)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

print("Classification Report:\n", classification_report(y_test, y_pred_pca))

```

Number of PCA components selected: 36

Naive Bayes Accuracy (PCA Features): 0.7700



Classification Report:

	precision	recall	f1-score	support
1.0	0.97	0.80	0.87	1172
2.0	0.87	0.35	0.50	94

3.0	0.17	0.33	0.22	3
4.0	0.19	0.97	0.32	64
5.0	0.67	0.67	0.67	3
6.0	0.33	0.50	0.40	2
8.0	0.00	0.00	0.00	1
accuracy			0.77	1339
macro avg	0.46	0.52	0.43	1339
weighted avg	0.92	0.77	0.82	1339

```

/home/austi/miniconda3/envs/msse/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/austi/miniconda3/envs/msse/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/austi/miniconda3/envs/msse/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```