

Sam_failed_life_models

May 13, 2025

0.1 Unsupervised Learning to Predict Race from Merged Dataset

[]:

```
[23]: #imports
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

0.2 Load Data and Inspect

```
[24]: full_data_df = pd.read_csv('final_merged.csv')
full_data_df

print("Shape of the data: ", full_data_df.shape)

full_data_df.info()
full_data_df.head()
```

Shape of the data: (9222, 135)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9222 entries, 0 to 9221

Columns: 135 entries, Patient ID to Survival Months

dtypes: float64(74), int64(2), object(59)

memory usage: 9.5+ MB

C:\Users\smcca\AppData\Local\Temp\ipykernel_48032\3584519652.py:1: DtypeWarning:

Columns (5,7,12,15,16,17,18,30,65,82,83,91,92,94,95,96,97,98,99,100,101,102,103,104,105,106,107) have mixed types. Specify dtype option on import or set

low_memory=False.

```
full_data_df = pd.read_csv('final_merged.csv')
```

```
[24]: Patient ID  Age  Surgery  Cancer Type  Cancer Type Detailed \
0    MB-0110   45      1.0   Breast Cancer  Invasive Breast Carcinoma
1    MB-0156   60      1.0   Breast Cancer  Invasive Breast Carcinoma
2    MB-0159   51      1.0   Breast Cancer  Invasive Breast Carcinoma
```

3	MB-0230	41	1.0	Breast Cancer	Breast Invasive Ductal Carcinoma
4	MB-0281	52	1.0	Breast Sarcoma	Breast Angiosarcoma

	Cellularity	Chemotherapy	Pam50 + Claudin-low subtype	Cohort	\
0	Low	0.0	claudin-low	1.0	
1	Low	0.0	claudin-low	1.0	
2	Low	0.0	Basal	1.0	
3	NaN	0.0	Her2	1.0	
4	NaN	0.0	claudin-low	1.0	

	ER status measured by IHC	...	Marital Status	T Stage	6th Stage	Grade	\
0	NaN	...	NaN	NaN	NaN	NaN	
1	NaN	...	NaN	NaN	NaN	NaN	
2	NaN	...	NaN	NaN	NaN	NaN	
3	NaN	...	NaN	NaN	NaN	NaN	
4	NaN	...	NaN	NaN	NaN	NaN	

	A Stage	Estrogen Status	Progesterone Status	Regional Node Examined	\
0	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	

	Reginol Node Positive	Survival Months
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

[5 rows x 135 columns]

```
[3]: 'Radio Therapy' in full_data_df.columns
```

```
[3]: True
```

```
[4]: 'Clinical Response, Evaluated Through Imaging ' in full_data_df.columns
```

```
[4]: True
```

```
[5]: full_data_df['Radio Therapy'].value_counts()
```

```
[5]: Radio Therapy
0.0    5702
1.0    1233
Name: count, dtype: int64
```

```
[6]: full_data_df['Chemotherapy'].value_counts()
```

```
[6]: Chemotherapy
0.0    4613
1.0    2252
Name: count, dtype: int64
```

```
[7]: full_data_df['T Stage'].value_counts()
```

```
[7]: T Stage
2.0    4053
1.0    3484
3.0     802
4.0     132
0.0      24
Name: count, dtype: int64
```

```
[8]: full_data_df['Survival Months'].value_counts()
```

```
[8]: Survival Months
69.0     119
63.0     116
66.0     115
102.0    114
50.0     109
...
7.0        5
26.0       5
11.0       4
2.0        4
1.0        1
Name: count, Length: 107, dtype: int64
```

```
[9]: full_data_df['N Stage'].value_counts()
```

```
[9]: N Stage
1.0    4779
2.0     874
0.0     529
3.0     507
Name: count, dtype: int64
```

```
[25]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import numpy as np

# Define target
target_col = 'Survival Months'

# Select numeric features
num_features = full_data_df.select_dtypes(include=[np.number])
num_features = num_features.drop(columns=[target_col])
```

```
[26]: # Define the threshold (60% of the total number of rows)
threshold = 0.2 * len(full_data_df)

# Drop columns where the number of NaN values is greater than the threshold
full_data_df_cleaned = full_data_df.dropna(thresh=threshold, axis=1)
```

```
[27]: num_features
```

```
[27]:
```

	Age	Surgery	Chemotherapy	Cohort	ER	Neoplasm	Histologic	Grade	\
0	45	1.0	0.0	1.0	0.0			NaN	
1	60	1.0	0.0	1.0	0.0			1.0	
2	51	1.0	0.0	1.0	0.0			NaN	
3	41	1.0	0.0	1.0	0.0			NaN	
4	52	1.0	0.0	1.0	0.0			NaN	
...		
9217	43	0.0	NaN	1.0	0.0			3.0	
9218	59	0.0	NaN	1.0	0.0			2.0	
9219	64	0.0	NaN	1.0	0.0			2.0	
9220	72	0.0	NaN	9.0	0.0			3.0	
9221	49	NaN	NaN	NaN	1.0			NaN	

	HER2	Histologic type	Lymph nodes examined	positive	Mutation	Count	\
0	0.0	0			NaN	NaN	
1	0.0	0			NaN	NaN	
2	0.0	0			NaN	NaN	
3	0.0	0			0.0	2.0	
4	0.0	0			NaN	NaN	
...		
9217	0.0	0			2.0	5.0	
9218	0.0	0			2.0	4.0	
9219	0.0	0			1.0	3.0	
9220	0.0	0			0.0	NaN	
9221	0.0	0			NaN	NaN	

	...	\
0	...	
1	...	
2	...	

3 ...
 4 ...

 9217 ...
 9218 ...
 9219 ...
 9220 ...
 9221 ...

Pathologic response to Neoadjuvant therapy: Pathologic stage (T) following
 neoadjuvant therapy \

0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
9217	NaN
9218	NaN
9219	NaN
9220	NaN
9221	NaN

Pathologic response to Neoadjuvant therapy: Pathologic stage (N)
 following neoadjuvant therapy \

0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
9217	NaN
9218	NaN
9219	NaN
9220	NaN
9221	NaN

Pathologic response to Neoadjuvant therapy: Pathologic stage (M)
 following neoadjuvant therapy \

0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
9217	NaN
9218	NaN

9219	NaN
9220	NaN
9221	NaN

Overall Near-complete Response: Stricter Definition \	
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
9217	NaN
9218	NaN
9219	NaN
9220	NaN
9221	NaN

Overall Near-complete Response: Looser Definition \	
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN
...	...
9217	NaN
9218	NaN
9219	NaN
9220	NaN
9221	NaN

Near-complete Response (Graded Measure)	Race	N	Stage \
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN
...	
9217	NaN	NaN	NaN
9218	NaN	NaN	NaN
9219	NaN	NaN	NaN
9220	NaN	NaN	NaN
9221	NaN	1.0	NaN

Regional Node Examined	Reginol Node Positive
0	NaN
1	NaN
2	NaN

3	NaN	NaN
4	NaN	NaN
...
9217	NaN	NaN
9218	NaN	NaN
9219	NaN	NaN
9220	NaN	NaN
9221	NaN	NaN

[9222 rows x 75 columns]

[]:

[28]: `len(num_features.columns)`

[28]: 75

```
[29]: # Ensure num_features only contains numeric columns

from sklearn.impute import SimpleImputer
# Initialize the imputer to fill NaNs with the mean
imputer = SimpleImputer(strategy='mean')

# Ensure that num_features only contains numeric columns
num_features = full_data_df.select_dtypes(include=[np.number])

# Check the shape of num_features to make sure it has the expected number of
↳ columns
print("Shape of num_features:", num_features.shape)

# List of columns in num_features
print("Columns in num_features:", num_features.columns)
print("Shape of num_features:", num_features.shape)

# Apply the imputer to your DataFrame
full_data_df_imputed = pd.DataFrame(imputer.
↳ fit_transform(num_features), columns=num_features.columns)

# Check the result
print(full_data_df_imputed.isna().sum()) # Should show 0 NaN values
```

Shape of num_features: (9222, 76)

Columns in num_features: Index(['Age', 'Surgery', 'Chemotherapy', 'Cohort',
'ER',

'Neoplasm Histologic Grade', 'HER2', 'Histologic type',

'Lymph nodes examined positive', 'Mutation Count',

```

'Nottingham prognostic index', 'Overall Survival (Months)', 'PR',
'Radio Therapy', 'Relapse Free Status (Months)', 'Tumor Size',
'T Stage', 'Days to MRI (From the Date of Diagnosis)', 'Manufacturer',
'Manufacturer Model Name', 'Scan Options', 'Field Strength (Tesla)',
'Patient Position During MRI', 'Contrast Agent',
'Contrast Bolus Volume (mL)', 'TR (Repetition Time)', 'TE (Echo Time)',
'Acquisition Matrix', 'Slice Thickness ', 'Rows', 'Columns',
'Reconstruction Diameter ', 'Flip Angle ',
'FOV Computed (Field of View) in cm ', 'Date of Birth (Days)',
'Menopause (at diagnosis)', 'Race and Ethnicity',
'Metastatic at Presentation (Outside of Lymph Nodes)', 'Mol Subtype',
'Oncotype score', 'Staging(Tumor Size)# [T]',
'Staging(Nodes)#(Nx replaced by -1)[N]',
'Staging(Metastasis)#(Mx -replaced by -1)[M]',
'Tumor Grade(T) (Tubule)', 'Tumor Grade(N)(Nuclear)',
'Tumor Grade(M)(Mitotic)', 'Nottingham grade',
'Multicentric/Multifocal', 'Contralateral Breast Involvement',
'Lymphadenopathy or Suspicious Nodes', 'Skin/Nipple Involvement',
'Pec/Chest Involvement', 'Adjuvant Radiation Therapy',
'Clinical Response, Evaluated Through Imaging ',
'Pathologic Response to Neoadjuvant Therapy', 'Recurrence event(s)',
'Age at last contact in EMR f/u(days)(from the date of diagnosis) ,last
time patient known to be alive, unless age of death is reported(in such case the
age of death',
'Adjuvant Chemotherapy', 'Neoadjuvant Endocrine Therapy Medications ',
'Adjuvant Endocrine Therapy Medications ', 'Known Ovarian Status ',
'Therapeutic or Prophylactic Oophorectomy as part of Endocrine Therapy ',
'Neoadjuvant Anti-Her2 Neu Therapy', 'Adjuvant Anti-Her2 Neu Therapy ',
'Received Neoadjuvant Therapy or Not',
'Pathologic response to Neoadjuvant therapy: Pathologic stage (T)
following neoadjuvant therapy ',
'Pathologic response to Neoadjuvant therapy: Pathologic stage (N)
following neoadjuvant therapy',
'Pathologic response to Neoadjuvant therapy: Pathologic stage (M)
following neoadjuvant therapy ',
'Overall Near-complete Response: Stricter Definition',
'Overall Near-complete Response: Looser Definition',
'Near-complete Response (Graded Measure)', 'Race', 'N Stage',
'Regional Node Examined', 'Reginol Node Positive', 'Survival Months'],
dtype='object')
Shape of num_features: (9222, 76)
Age 0
Surgery 0
Chemotherapy 0
Cohort 0
ER 0
..
Race 0

```



```
N Stage          0
Regional Node Examined  0
Reginol Node Positive  0
Survival Months      0
Length: 76, dtype: int64
```

```
[31]: target_col = 'Survival Months'

# Define features and target from the imputed DataFrame
X = full_data_df_imputed # Features from imputed data
y = full_data_df_imputed[target_col] # Target from imputed data

# Drop the target column from X (features)
X = X.drop(columns=[target_col])

# Drop NaNs from features and align target
X = X.dropna()
y = y.loc[X.index]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

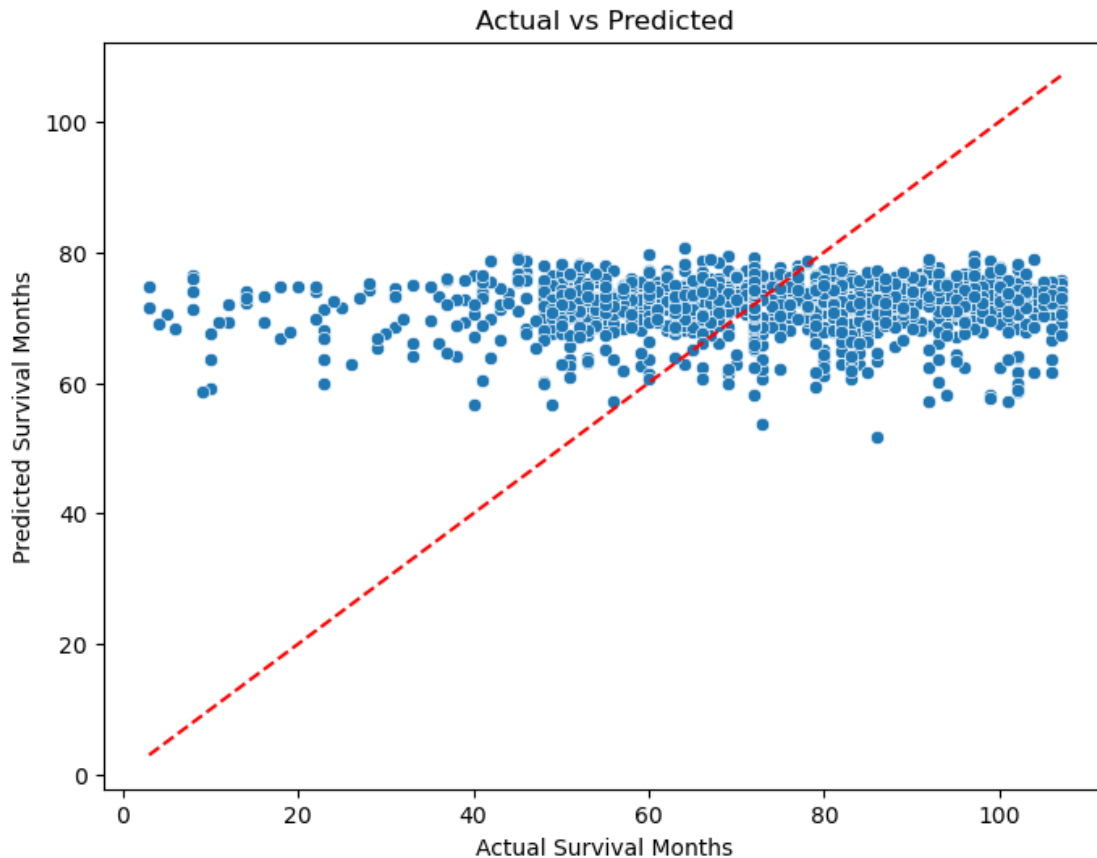
# Fit linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

# Evaluate
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))

# Optional: plot predicted vs actual
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred)
plt.xlabel("Actual Survival Months")
plt.ylabel("Predicted Survival Months")
plt.title("Actual vs Predicted")
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
    color='red', linestyle='--')
plt.show()
```

```
Mean Squared Error: 307.16179089356694
R^2 Score: -0.018432886493055856
```



```
[20]: import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

model = LinearRegression()
model.fit(X_train, y_train)

# Prediction
y_pred = model.predict(X_test)

# Evaluation
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))
print("Intercept:", model.intercept_)
print("Coefficient:", model.coef_[0])
```

```
-----
ValueError                                Traceback (most recent call last)
~\AppData\Local\Temp\ipykernel_19180\1385663674.py in ?()
```

```

3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import mean_squared_error, r2_score
5
6 model = LinearRegression()
----> 7 model.fit(X_train, y_train)
8
9 # Prediction
10 y_pred = model.predict(X_test)

c:\Users\smcca\anaconda3\Lib\site-packages\sklearn\base.py in ?(estimator,
↳ *args, **kwargs)
1469         skip_parameter_validation=(
1470             prefer_skip_nested_validation or
↳ global_skip_validation
1471         )
1472     ):
-> 1473         return fit_method(estimator, *args, **kwargs)

c:\Users\smcca\anaconda3\Lib\site-packages\sklearn\linear_model\_base.py in ?
↳ (self, X, y, sample_weight)
605         n_jobs_ = self.n_jobs
606
607         accept_sparse = False if self.positive else ["csr", "csc", "coo ]
608
--> 609         X, y = self._validate_data(
610             X,
611             y,
612             accept_sparse=accept_sparse,

c:\Users\smcca\anaconda3\Lib\site-packages\sklearn\base.py in ?(self, X, y,
↳ reset, validate_separately, cast_to_ndarray, **check_params)
646         if "estimator" not in check_y_params:
647             check_y_params = {**default_check_params,
↳ **check_y_params}
648         y = check_array(y, input_name="y", **check_y_params)
649     else:
--> 650         X, y = check_X_y(X, y, **check_params)
651         out = X, y
652
653         if not no_val_X and check_params.get("ensure_2d", True):

c:\Users\smcca\anaconda3\Lib\site-packages\sklearn\utils\validation.py in ?(X,
↳ y, accept_sparse, accept_large_sparse, dtype, order, copy, force_writeable,
↳ force_all_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples,
↳ ensure_min_features, y_numeric, estimator)
1297         raise ValueError(
1298             f"{estimator_name} requires y to be passed, but the target :
↳ is None"

```

```

1299         )
1300
-> 1301     X = check_array(
1302         X,
1303         accept_sparse=accept_sparse,
1304         accept_large_sparse=accept_large_sparse,

c:\Users\smcca\anaconda3\Lib\site-packages\sklearn\utils\validation.py in ?
-> (array, accept_sparse, accept_large_sparse, dtype, order, copy,
-> force_writeable, force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
-> ensure_min_features, estimator, input_name)
1009         )
1010         array = xp.astype(array, dtype, copy=False)
1011     else:
1012         array = _asarray_with_order(array, order=order,
-> dtype=dtype, xp=xp)
-> 1013     except ComplexWarning as complex_warning:
1014         raise ValueError(
1015             "Complex data not supported\n{}\n".format(array)
1016         ) from complex_warning

c:\Users\smcca\anaconda3\Lib\site-packages\sklearn\utils\_array_api.py in ?
-> (array, dtype, order, copy, xp, device)
747     # Use NumPy API to support order
748     if copy is True:
749         array = numpy.array(array, order=order, dtype=dtype)
750     else:
--> 751         array = numpy.asarray(array, order=order, dtype=dtype)
752
753     # At this point array is a NumPy ndarray. We convert it to an
-> array
754     # container that is consistent with the input's namespace.

c:\Users\smcca\anaconda3\Lib\site-packages\pandas\core\generic.py in ?(self,
-> dtype, copy)
2149     def __array__(
2150         self, dtype: npt.DTypeLike | None = None, copy: bool_t | None =
-> None
2151     ) -> np.ndarray:
2152         values = self._values
-> 2153         arr = np.asarray(values, dtype=dtype)
2154         if (
2155             astype_is_view(values.dtype, arr.dtype)
2156             and using_copy_on_write()

```

ValueError: could not convert string to float: 'Positive'

0.3 Data Cleaning

```
[15]: #Dropping any duplicate rows
full_data_df.drop_duplicates(inplace=True)

print('Radio Therapy' in full_data_df.columns)

# Drop features with a lot of missing values
missing_ratio = full_data_df.isnull().mean()
columns_to_drop = missing_ratio[missing_ratio > 0.9].index
full_data_df.drop(columns=columns_to_drop, inplace=True)

print('Radio Therapy' in full_data_df.columns)
```

True

True

```
[16]: # Look at new shape of the data
print("Shape of the data after dropping columns with a lot of missing values:␣
↪", full_data_df.shape)
```

Shape of the data after dropping columns with a lot of missing values: (9221, 135)

```
[ ]:
```

```
[17]: #inspecting the race column
full_data_df['Overall Near-complete Response: Stricter Definition'].
↪value_counts(dropna=False)

#Drop rows with missing values in Race column

full_data_df = full_data_df.dropna(subset=['Overall Near-complete Response: ␣
↪Stricter Definition'])

#Drop rows where race is = 0
full_data_df = full_data_df[full_data_df['Overall Near-complete Response: ␣
↪Stricter Definition'] != 0]

full_data_df['Overall Near-complete Response: Stricter Definition'].
↪value_counts()

full_data_df.drop('Clinical Response, Evaluated Through Imaging ',axis=1)

full_data_df.drop('Pathologic Response to Neoadjuvant Therapy',axis=1)
```

```

full_data_df.drop('Pathologic Response to Neoadjuvant Therapy',axis=1)

full_data_df.drop('Overall Near-complete Response: Looser Definition',axis=1)

full_data_df.drop('Near-complete Response (Graded Measure)',axis=1)

full_data_df['Radio Therapy']

```

```

[17]: 2739    0.0
      3731    0.0
      3737    0.0
      3738    0.0
      3739    0.0
      ...
      5627    0.0
      6816    1.0
      6817    1.0
      6818    1.0
      6934    1.0
Name: Radio Therapy, Length: 572, dtype: float64

```

```

[ ]: from sklearn.model_selection import train_test_split
     from sklearn.naive_bayes import GaussianNB
     from sklearn.metrics import accuracy_score, confusion_matrix
     import matplotlib.pyplot as plt
     import seaborn as sns

     # Define target
     target_col = 'Overall Near-complete Response: Stricter Definition'

     # Define features and target
     X = full_data_df[['Radio Therapy', 'Chemotherapy', 'Tumor Size', 'Age', 'T_
     ↳Stage', 'Survival Months']]
     y = full_data_df[target_col]

     # Drop NaNs from features and align target
     X = X.dropna()
     y = y.loc[X.index]

     # Drop classes with fewer than 2 samples
     class_counts = y.value_counts()
     valid_classes = class_counts[class_counts >= 2].index
     X = X[y.isin(valid_classes)]

```

```

y = y[y.isin(valid_classes)]

# Re-align X and y again just to be sure
X = X.loc[y.index]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, stratify=y, test_size=0.2, random_state=42
)

# Ready to test different subset sizes

```

[19]: X_train

```

[19]:      Radio Therapy  Chemotherapy  Tumor Size  Age  T Stage  Survival Months
4055              0.0              1.0        17.0   49        1.0             50.0
3739              0.0              1.0        10.0   60        1.0             51.0
4249              0.0              1.0        16.0   40        1.0             66.0
4199              0.0              1.0        10.0   47        1.0             85.0
5573              0.0              1.0        69.0   54        3.0             74.0
...              ...              ...        ...   ...        ...             ...
4216              0.0              1.0        16.0   47        1.0             76.0
4298              0.0              1.0        20.0   59        1.0             99.0
4449              0.0              1.0        40.0   51        2.0             67.0
4749              0.0              1.0        34.0   59        2.0             81.0
4761              0.0              1.0        22.0   53        2.0             61.0

```

[402 rows x 6 columns]

```

[27]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Define target
target_col = 'Overall Near-complete Response: Stricter Definition'

# Define features and target
X = full_data_df[['Radio Therapy', 'Chemotherapy', 'Tumor Size', 'Age', 'T_
↳Stage', 'Survival Months']]
y = full_data_df[target_col]

# Drop NaNs from features and align target
X = X.dropna()
y = y.loc[X.index]

```

```

# Drop classes with fewer than 2 samples
class_counts = y.value_counts()
valid_classes = class_counts[class_counts >= 2].index
X = X[y.isin(valid_classes)]
y = y[y.isin(valid_classes)]

# Re-align X and y again just to be sure
X = X.loc[y.index]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, stratify=y, test_size=0.2, random_state=42
)

# Ready to test different subset sizes

```

```
[10]: results = {}
```

```

[26]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
import random

selected_features = ['Radio Therapy', 'Chemotherapy', 'Age', 'T Stage']

model = GaussianNB()
model.fit(X_train[selected_features], y_train)
y_pred = model.predict(X_test[selected_features])

acc = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)

size = 4

results[size] = {
    "features": selected_features,
    "accuracy": acc,
    "conf_matrix": conf_matrix
}

print(f"\nFeature Subset Size: {size}")
print("Selected Features:", selected_features)
print(f"Accuracy: {acc:.4f}")

plt.figure(figsize=(6, 4))

```




```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title(f"Confusion Matrix (Feature Size: {size})")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

Feature Subset Size: 4

Selected Features: ['Radio Therapy', 'Chemotherapy', 'Age', 'T Stage']

Accuracy: 0.7030



[]:

0.4 EDA

[12]: full_data_df.describe()

```
[12]:
```

	Age	Surgery	Chemotherapy	Cohort	ER \
count	572.000000	27.0	572.000000	0.0	572.000000
mean	50.846154	0.0	0.998252	NaN	0.603147
std	7.900039	0.0	0.041812	NaN	0.489673
min	21.000000	0.0	0.000000	NaN	0.000000
25%	47.000000	0.0	1.000000	NaN	0.000000

50%	53.000000	0.0	1.000000	NaN	1.000000
75%	58.000000	0.0	1.000000	NaN	1.000000
max	73.000000	0.0	1.000000	NaN	1.000000

	Neoplasms Histologic Grade	HER2	Histologic type \
count	0.0	572.000000	572.000000
mean	NaN	0.477273	0.669580
std	NaN	0.499920	0.496133
min	NaN	0.000000	0.000000
25%	NaN	0.000000	0.000000
50%	NaN	0.000000	1.000000
75%	NaN	1.000000	1.000000
max	NaN	1.000000	2.000000

	Lymph nodes examined positive	Mutation Count ... \
count	0.0	0.0 ...
mean	NaN	NaN ...
std	NaN	NaN ...
min	NaN	NaN ...
25%	NaN	NaN ...
50%	NaN	NaN ...
75%	NaN	NaN ...
max	NaN	NaN ...

Pathologic response to Neoadjuvant therapy: Pathologic stage (N)
following neoadjuvant therapy \

count	378.000000
mean	0.142857
std	0.515761
min	-1.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

Pathologic response to Neoadjuvant therapy: Pathologic stage (M)
following neoadjuvant therapy \

count	376.000000
mean	-0.997340
std	0.051571
min	-1.000000
25%	-1.000000
50%	-1.000000
75%	-1.000000
max	0.000000

Overall Near-complete Response: Stricter Definition \

count	572.000000
mean	1.503497
std	0.649657
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	3.000000

	Overall Near-complete Response: Looser Definition \
count	572.000000
mean	1.503497
std	0.649657
min	1.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	3.000000

	Near-complete Response (Graded Measure)	Race	N Stage \
count	572.000000	572.000000	568.000000
mean	3.524476	1.152098	1.010563
std	5.058789	0.631579	0.432247
min	1.000000	0.000000	0.000000
25%	1.000000	1.000000	1.000000
50%	1.000000	1.000000	1.000000
75%	4.000000	1.000000	1.000000
max	19.000000	7.000000	3.000000

	Regional Node Examined	Reginol Node Positive	Survival Months
count	503.000000	503.000000	503.000000
mean	12.789264	2.075547	71.854871
std	8.106815	2.453199	22.094888
min	1.000000	1.000000	4.000000
25%	7.000000	1.000000	58.000000
50%	12.000000	1.000000	73.000000
75%	18.000000	2.000000	89.000000
max	49.000000	26.000000	107.000000


[8 rows x 76 columns]

[]:

[13]: *#Coorelation heatmap for numerical features*

```
numeric_features = full_data_df.select_dtypes(include=[np.number])
plt.figure(figsize=(12, 8))
```

```
sns.heatmap(numeric_features.corr(), annot=False, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



0.5 Guassian and PCA

```
[14]: from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np

# Define target
target_col = 'Overall Near-complete Response: Stricter Definition'

# Drop columns with too many NaNs (>80%)
full_data_df = full_data_df.loc[:, full_data_df.isnull().mean() < 0.8]
```

```

# Keep only numeric columns and exclude the target
numeric_cols = full_data_df.select_dtypes(include=[np.number]).columns
features = numeric_cols.drop(target_col, errors='ignore') # in case target_col
↳ isn't numeric

# Strategy 1: Drop rows with any NaNs
df_dropped = full_data_df[features.union([target_col])].dropna()

# Strategy 2: Fill numeric NaNs with mean
df_filled = full_data_df[features.union([target_col])].copy()
df_filled[features] = df_filled[features].fillna(df_filled[features].mean())

# Scale features only if rows exist
scaler = StandardScaler()

if not df_dropped.empty:
    df_dropped[features] = scaler.fit_transform(df_dropped[features])
else:
    print("No rows left in df_dropped after dropna().")

if not df_filled.empty:
    df_filled[features] = scaler.fit_transform(df_filled[features])
else:
    print("No rows left in df_filled after fillna().")

print(" Dropped dataset shape:", df_dropped.shape)
print(" Filled dataset shape:", df_filled.shape)

```

No rows left in df_dropped after dropna().
Dropped dataset shape: (0, 67)
Filled dataset shape: (572, 67)

```

[15]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
import random

# Define features and target
X = df_filled[features]
y = df_filled[target_col]

# Drop classes with fewer than 2 samples
class_counts = y.value_counts()
valid_classes = class_counts[class_counts >= 2].index

```

```

X = X[y.isin(valid_classes)]
y = y[y.isin(valid_classes)]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
↳test_size=0.2, random_state=42)

# Test different subset sizes
feature_sizes = [3, 5, 7, 10]
results = {}

for size in feature_sizes:
    selected_features = random.sample(list(X.columns), size)

    model = GaussianNB()
    model.fit(X_train[selected_features], y_train)
    y_pred = model.predict(X_test[selected_features])

    acc = accuracy_score(y_test, y_pred)
    conf_matrix = confusion_matrix(y_test, y_pred)

    results[size] = {
        "features": selected_features,
        "accuracy": acc,
        "conf_matrix": conf_matrix
    }

print(f"\nFeature Subset Size: {size}")
print("Selected Features:", selected_features)
print(f"Accuracy: {acc:.4f}")


plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.title(f"Confusion Matrix (Feature Size: {size})")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

```

Feature Subset Size: 3
Selected Features: ['Reginol Node Positive', 'Contrast Bolus Volume (mL)',
'Known Ovarian Status ']
Accuracy: 0.6174


```



Feature Subset Size: 5

Selected Features: ['Recurrence event(s)', 'Adjuvant Radiation Therapy', 'Pathologic response to Neoadjuvant therapy: Pathologic stage (T) following neoadjuvant therapy ', 'Mol Subtype', 'Race']


Accuracy: 0.5043



Feature Subset Size: 7


Selected Features: ['Pec/Chest Involvement', 'Skin/Nipple Invovlement', 'Columns', 'Chemotherapy', 'Race', 'Slice Thickness ', 'Manufacturer']

Accuracy: 0.5391



Feature Subset Size: 10

Selected Features: ['Survival Months', 'Neoadjuvant Anti-Her2 Neu Therapy', 'Pathologic response to Neoadjuvant therapy: Pathologic stage (N) following neoadjuvant therapy', 'Reconstruction Diameter ', 'Staging(Tumor Size)# [T]', 'Scan Options', 'Tumor Grade(T) (Tubule)', 'Adjuvant Endocrine Therapy Medications ', 'Days to MRI (From the Date of Diagnosis)', 'Chemotherapy']
Accuracy: 0.6957



[]:

[]:

```
[16]: # Baseline model
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Train full-feature Naive Bayes model
model_full = GaussianNB()
model_full.fit(X_train, y_train)

# Predict on test set
y_pred_full = model_full.predict(X_test)

# Evaluate
accuracy_full = accuracy_score(y_test, y_pred_full)
print(f"Naive Bayes Accuracy (All Features): {accuracy_full:.4f}")

# Classification report
```

```

print("Classification Report:\n", classification_report(y_test, y_pred_full))

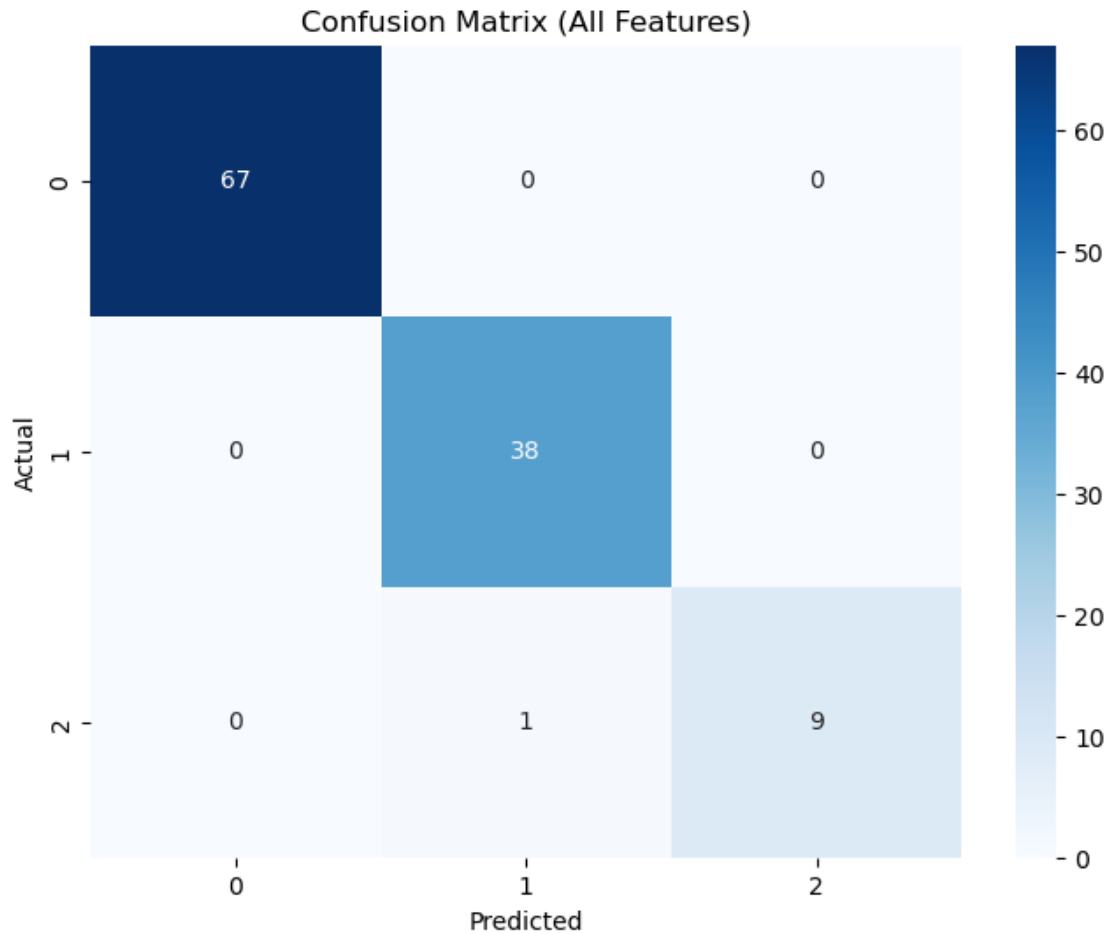
# Confusion Matrix
conf_matrix_full = confusion_matrix(y_test, y_pred_full)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_full, annot=True, fmt="d", cmap="Blues")
plt.title("Confusion Matrix (All Features)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```

Naive Bayes Accuracy (All Features): 0.9913

Classification Report:


	precision	recall	f1-score	support
1.0	1.00	1.00	1.00	67
2.0	0.97	1.00	0.99	38
3.0	1.00	0.90	0.95	10
accuracy			0.99	115
macro avg	0.99	0.97	0.98	115
weighted avg	0.99	0.99	0.99	115



```
[17]: from sklearn.decomposition import PCA

# Reduce to 2D for visualization
pca_vis = PCA(n_components=2)
X_train_pca_2d = pca_vis.fit_transform(X_train)

# Plot PCA components colored by Race
plt.figure(figsize=(8, 6))
sns.scatterplot(x=X_train_pca_2d[:, 0], y=X_train_pca_2d[:, 1], hue=y_train,
               palette="Set1")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.title("2D PCA Projection Colored by Race")
plt.legend(title="Race")
plt.grid(True)
plt.tight_layout()
plt.show()
```



0.6 Evaluation

```
[18]: from sklearn.decomposition import PCA
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report

# Fit PCA to retain 90% variance
pca_final = PCA(n_components=0.90)
X_train_pca = pca_final.fit_transform(X_train)
X_test_pca = pca_final.transform(X_test)

print(f"Number of PCA components selected: {pca_final.n_components_}")

# Train and predict
model_pca = GaussianNB()
model_pca.fit(X_train_pca, y_train)
y_pred_pca = model_pca.predict(X_test_pca)

# Evaluate
```

```

acc_pca = accuracy_score(y_test, y_pred_pca)
print(f"Naive Bayes Accuracy (PCA Features): {acc_pca:.4f}")


conf_matrix_pca = confusion_matrix(y_test, y_pred_pca)
sns.heatmap(conf_matrix_pca, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix (PCA Features)")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

print("Classification Report:\n", classification_report(y_test, y_pred_pca))

```

Number of PCA components selected: 23

Naive Bayes Accuracy (PCA Features): 0.9565



Classification Report:

	precision	recall	f1-score	support
1.0	0.97	0.96	0.96	67
2.0	0.92	0.95	0.94	38

3.0	1.00	1.00	1.00	10
accuracy			0.96	115
macro avg	0.96	0.97	0.97	115
weighted avg	0.96	0.96	0.96	115